

# Use Case UC1: Retrieve Relevant Document Chunks

**Scope:** RAG Chatbot System (Backend Pipeline)

**Level:** Sub-function

**Primary Actor:** User (via System Trigger)

## 1. Stakeholders and Interests

- **User:** Wants to find text segments that answer their question accurately. Expects fast and relevant results.
- **Developer:** Wants to ensure correct query cleaning, effective STOP\_WORDS filtering, and accurate ranking based on Term Frequency (TF).
- **Orchestrator (System Controller):** Expects the module to return data in a standard format (List<Hit>) without errors to proceed to the answering stage.

## 2. Preconditions

- The system is initialized, and the KeywordIndex is successfully loaded into memory.
- The user's question (question) has been received, and the user's intent (Intent) has been determined by the previous module (e.g., STAFF\_LOOKUP).

## 3. Success Guarantee (Postconditions)

- The system identifies the top \$K\$ (default 10) most relevant document chunks (Hit) based on the query, sorts them according to the scoring rules, and saves them into the Context object.

## 4. Main Success Scenario (Basic Flow)

1. The **Orchestrator** initiates the **QueryWriter** strategy with the user's question and the detected intent.
2. The **QueryWriter** normalizes the question by converting it to lowercase and removing punctuation marks.
3. The **QueryWriter** iterates through the tokens and filters out meaningless conjunctions and banned words defined in STOP\_WORDS.
4. The **QueryWriter** checks the Intent and appends specific booster terms to the list (e.g., adding "office", "email" for STAFF\_LOOKUP intent).
5. The **QueryWriter** returns the list of cleaned and expanded terms (List<String>) to the Orchestrator.
6. The **Orchestrator** calls the **Retriever** strategy using these terms and the loaded KeywordIndex.
7. The **Retriever** fetches matching records (IndexEntry) from the KeywordIndex for each query term.

8. The **Retriever** calculates a cumulative score for each document chunk based on the sum of Term Frequencies (TF).
9. The **Retriever** sorts the results based on the following deterministic tie-breaking rules:
  - o Score (Descending order)
  - o Document ID (Ascending order - if scores are equal)
  - o Chunk ID (Ascending order - if scores and document IDs are equal)
10. The **Retriever** returns the top 10 highest-scoring results (List<Hit>) to the Orchestrator.

## 5. Extensions (Alternative Flows)

- **3a. The query consists entirely of "Stop Words" (e.g., "and or but"):**
  - o **3a1.** The QueryWriter detects that the filtered list is empty.
  - o **3a2.** The system uses the original words from the question (unfiltered) or falls back to default general terms.
- **7a. No matches found in the Index for the query terms:**
  - o **7a1.** The Retriever creates an empty list of Hits.
  - o **7a2.** The system logs a warning: "No hits found for terms: [terms]".
  - o **7a3.** Returns the empty list (The AnswerAgent will handle the "No answer found" response in the next stage).
- **10a. Fewer than 10 results found:**
  - o **10a1.** The Retriever returns all available results (e.g., 3 hits) without attempting to pad the list.

## 6. Special Requirements

- **Deterministic Behavior:** The same input (Question + Configuration) must always produce the exact same results in the exact same order. Tie-break rules must be strictly enforced.
- **Performance:** The retrieval process must complete in under 200ms, even with a fully loaded index.
- **Language Support:** Turkish and English characters (e.g., İ, ı, ş, ç) must be preserved during normalization.

## 7. Technology and Data Variations

- The index structure is read from a JSON file and stored in memory as a HashMap within the KeywordIndex class.
- Sorting algorithms are implemented using Java's built-in Comparator interface.
- Logging and tracing are handled via the TraceBus (Observer Pattern).